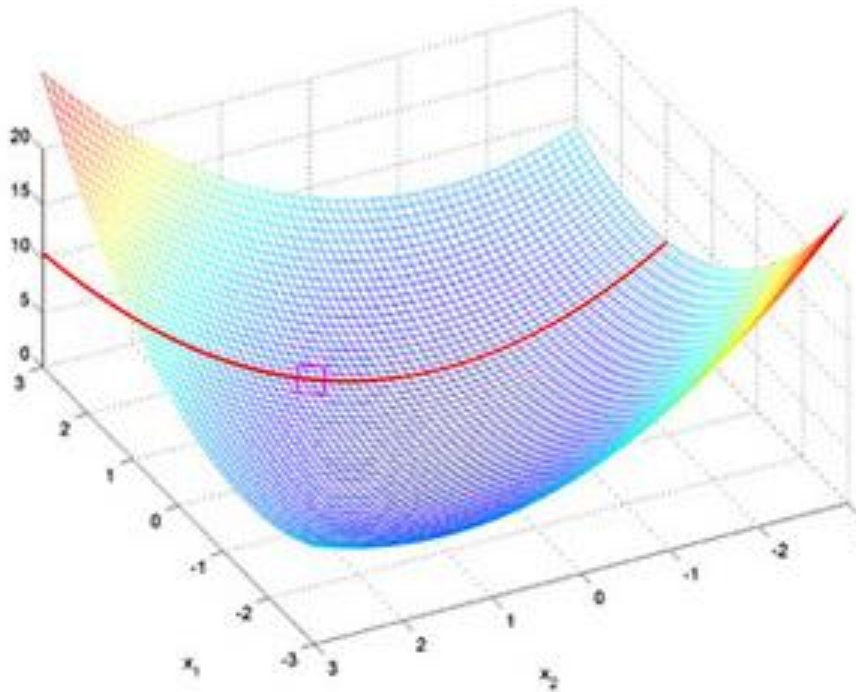


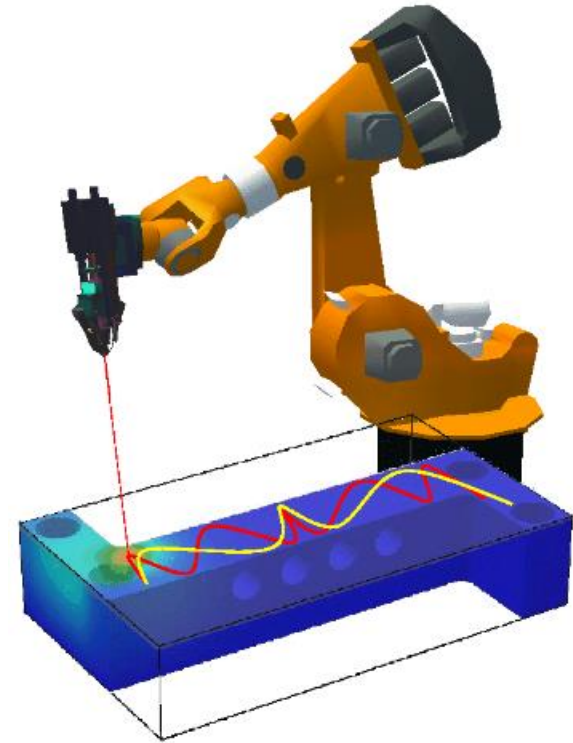
MEGR 7090/8090: Advanced Optimal Control



$$V_n(\mathbf{x}_n) = \min_{\{\mathbf{u}_n, \mathbf{u}_{n+1}, \dots, \mathbf{u}_{N-1}\}} \left[\frac{1}{2} \sum_{k=n}^{N-1} (\mathbf{x}_k^T \mathbf{Q}_k \mathbf{x}_k + \mathbf{u}_k^T \mathbf{R} \mathbf{u}_k) + \frac{1}{2} \mathbf{x}_N^T \mathbf{Q}_N \mathbf{x}_N \right]$$

$$\begin{aligned} V_n(\mathbf{x}_n) &= \min_{\{\mathbf{u}_n, \mathbf{u}_{n+1}, \dots, \mathbf{u}_{N-1}\}} \left[\frac{1}{2} \sum_{k=n}^{N-1} (\mathbf{x}_k^T \mathbf{Q}_k \mathbf{x}_k + \mathbf{u}_k^T \mathbf{R} \mathbf{u}_k) + \frac{1}{2} \mathbf{x}_N^T \mathbf{Q}_N \mathbf{x}_N \right] \\ &= \min_{\mathbf{u}_n} \left[\frac{1}{2} (\mathbf{x}_n^T \mathbf{Q}_n \mathbf{x}_n + \mathbf{u}_n^T \mathbf{R} \mathbf{u}_n) + \underbrace{\min_{\{\mathbf{u}_{n+1}, \dots, \mathbf{u}_{N-1}\}} \left[\frac{1}{2} \sum_{k=n+1}^{N-1} (\mathbf{x}_k^T \mathbf{Q}_k \mathbf{x}_k + \mathbf{u}_k^T \mathbf{R} \mathbf{u}_k) + \frac{1}{2} \mathbf{x}_N^T \mathbf{Q}_N \mathbf{x}_N \right]}_{V_{n+1}(\mathbf{x}_{n+1})} \right] \\ &= \min_{\mathbf{u}_n} \left[\frac{1}{2} (\mathbf{x}_n^T \mathbf{Q}_n \mathbf{x}_n + \mathbf{u}_n^T \mathbf{R} \mathbf{u}_n) + V_{n+1}(\mathbf{x}_{n+1}) \right] \end{aligned}$$

$$V_n(\mathbf{x}_n) = \min_{\mathbf{u}_n} \left[\frac{1}{2} (\mathbf{x}_n^T \mathbf{Q}_n \mathbf{x}_n + \mathbf{u}_n^T \mathbf{R} \mathbf{u}_n) + V_{n+1}(\mathbf{x}_{n+1}) \right]$$



Lecture 16
October 17, 2017

Dynamic Programming – Backward Recursion – Reminder



Getting set up:

- Quantize control variables into p discrete values
- Quantize state variables into q discrete values

Solution algorithm:

- Start at step $N-1$. For each of the q allowable state variables, calculate the stage cost $(g(\mathbf{x}(N-1), u(N-1)))$ for each of the p allowable control variables that lead to constraint satisfaction. Control variables that do not satisfy constraints are termed **inadmissible**. Record the optimal control signals and corresponding stage costs for each originating state.
- Move to step $N-2$. For each of the q allowable state variables, calculate the stage cost $(g(\mathbf{x}(N-1), u(N-1)))$ and associated intermediate state, $\mathbf{x}_i(N-1)$, for each of the p allowable control variables that lead to constraint satisfaction. To determine which control variable is optimal, compute the total cost to go as $J_{\mathbf{x}_o(N-2) \rightarrow \mathbf{x}_i(N-1) \rightarrow \mathbf{x}_f(N)}^* = J_{\mathbf{x}_o(N-2) \rightarrow \mathbf{x}_i(N-1)}^* + J_{\mathbf{x}_i(N-1) \rightarrow \mathbf{x}_f(N)}^*$
- Move to step $N-3$ and repeat the process (total cost to go is now $J_{\mathbf{x}_o(N-3) \rightarrow \mathbf{x}_i(N-2) \rightarrow \mathbf{x}_f(N)}^* = J_{\mathbf{x}_o(N-3) \rightarrow \mathbf{x}_i(N-2)}^* + J_{\mathbf{x}_i(N-2) \rightarrow \mathbf{x}_f(N)}^*$). Keep stepping backward in time until step 0.

Dynamic Programming – Forward Recursion – Reminder



Getting set up:

- Quantize control variables into p discrete values
- Quantize state variables into q discrete values
- Usually the initial state, $\mathbf{x}(0)$, is specified (you don't get to choose your initial state)

Solution algorithm:

- Start at step 1. For each of the q allowable values of $\mathbf{x}(1)$, back-compute the value of $\mathbf{x}(0)$ and stage cost that results from each of the p allowable control values. For each admissible control value (whose associated $\mathbf{x}(0)$ satisfies initial condition requirements), determine the optimal stage cost—this cost, denoted by $J_{\mathbf{x}_0(0) \rightarrow \mathbf{x}_t(1)}^*$, is the **optimal cost to arrive** at state $\mathbf{x}_t(1)$.
- Move to step 2. For each of the q allowable values of $\mathbf{x}(2)$, back-compute $\mathbf{x}(1)$ and the associated stage cost for each of the p allowable control variables that lead to constraint satisfaction. To determine which control sequence is optimal, compute the total cost to arrive as $J_{\mathbf{x}_0(0) \rightarrow \mathbf{x}_i(1) \rightarrow \mathbf{x}_f(2)}^* = J_{\mathbf{x}_0(0) \rightarrow \mathbf{x}_i(1)}^* + J_{\mathbf{x}_i(1) \rightarrow \mathbf{x}_f(2)}^*$
- Move to step 2 and repeat the process. Keep stepping forward in time until step N.

Dynamic Programming – Some Comments



Backward vs.
forward
recursion:

Why backward recursion is usually preferred?

1) Result of DP: u^* for every (quantized) x_0 .

- DP is really slow.

- Can implement a lookup table based on DP result.

- Can use result for benchmarking.

2) At step 1 of fwd. recursion, for every ^{quantized} state value, we compute x_0 for control value.

DP vs. SQP:

SQP: Local, continuous (u can take on any value)

DP: Global, gridded

Dynamic Programming – More Realistic Example



Suppose a vehicle's dynamics are given by:

$$m\dot{v} = u - C_{rr}mg - 0.5\rho v^2 C_d A_{ref}$$
$$\dot{x} = v$$

Parameter values: $m = 1000kg$, $g = 9.8 \frac{m}{s^2}$, $C_{rr} = 0.01$, $\rho = 1.2 \frac{kg}{m^3}$, $C_d = 0.4$, $A_{ref} = 5m^2$

Suppose that our goals are:

- Given an initial position of $x = 0$, achieve a final position of $x = 1600$ (approximately one “metric” mile of travel) after 60 seconds
- Minimize total energy expended over 60 seconds

Set up a nonlinear optimization problem and solve using dynamic programming

- Identify the states and control signal
- For different quantization levels, assess computational complexity

Dynamic Programming – More Realistic Example



States (2): Position (x) and velocity (v)

- Taking q_x and q_v as the number of quantized positions and velocities, respectively, the number of state variable combinations is given by $q = q_x q_v$

Control variable (1): Applied force (u)... Number of quantized forces denoted p

Horizon length: $N = \frac{60}{\Delta T}$ (where ΔT is the discrete time step)

Computational complexity: Number of cost function evaluations (E) given by:

$$E = Npq = Npq_x q_v$$

Example code available on Canvas.

Dynamic Programming – More Realistic Example

$$\text{Ex: } m\dot{v} = u - C_{rr}mg - 0.5\rho v^2 C_d A_{ref}$$

$$\dot{x} = v$$

$$v(k+1) = v(k) + \frac{\Delta t}{m} (u(k) - C_{rr}mg - 0.5\rho C_d A_{ref} v(k)^2)$$

$$x(k+1) = x(k) + \frac{v(k) + v(k+1)}{2} \Delta t \quad \left. \vphantom{x(k+1)} \right\} \text{trapezoidal approximation}$$

$$J(u; x(0)) = \int_0^{60} u(t) v(t) dt \approx \sum_{i=0}^{N-1} u(i) v(i) \Delta t$$

$$\text{Constraints: } x(N) \geq 1600$$

$$v(N) \geq 25$$


States : $\{x\} \leftarrow 21 \text{ values} \triangleq q_v$ $\{v\} \leftarrow 41 \text{ values} \triangleq q_x$

Control : $\{u\} \leftarrow 31 \text{ values} \triangleq p$

$$N=30 \Rightarrow \Delta t = 2 \text{ seconds.}$$

Dynamic Programming – More Realistic Example

Req'd # of ^{cost to go} computations: $N(q \times q \times p) = 800,730$
↑
curse of dimensionality



$$KE = \frac{1}{2}mv^2$$

Work must be done to increase kinetic energy. Therefore, the vehicle should finish at the minimum allowable speed. Note that its **average** speed must be greater than this to achieve $x(N) = 1600$ meters.

Dynamic Programming and the Principle of Optimality – A Famous Result



Consider the following system dynamics:

$$\mathbf{x}(k+1) = A\mathbf{x}(k) + Bu(k) \quad \mathbf{x} \in \mathbb{R}^n, u \in \mathbb{R}$$

...and the following objective function:

$$J(\mathbf{x}(0), \mathbf{u}) = \mathbf{x}^T(N)S\mathbf{x}(N) + \sum_{i=0}^{N-1} (\mathbf{x}^T(i)Q\mathbf{x}(i) + Ru^2(i))$$

where: $S \succ 0, Q \succ 0, R > 0$

- This is known as the **finite-horizon discrete-time linear quadratic regulator** problem
- Important notation (seems strange now, will make sense later): $P(0) \triangleq S$

Note: A detailed derivation for this problem is given in Kirk section 3.10

Discrete-Time LQR – Backward Recursion Step N-1

Note: $J_{\mathbf{x}(N) \rightarrow \mathbf{x}(N)}^* = \mathbf{x}^T(N)S\mathbf{x}(N) = \mathbf{x}^T(N)P(0)\mathbf{x}(N)$

Beginning backward recursion (step N-1):

$$\begin{aligned} J_{\mathbf{x}(N-1) \rightarrow \mathbf{x}(N)} &= \mathbf{x}^T(N-1)Q\mathbf{x}(N-1) + Ru^2(N-1) + \mathbf{x}^T(N)P(0)\mathbf{x}(N) \\ &= \mathbf{x}^T(N-1)Q\mathbf{x}(N-1) + Ru^2(N-1) + (A\mathbf{x}(N-1) + Bu(N-1))^T P(0)(A\mathbf{x}(N-1) + Bu(N-1)) \end{aligned}$$

Collecting like terms, differentiating with respect to $u(N-1)$, and setting the derivative to 0:

$$\frac{\partial J_{N-1 \rightarrow N}}{\partial u(N-1)} = 2(Ru(N-1) + B^T P(0)Bu(N-1) + B^T P(0)A\mathbf{x}(N-1))$$

$$u^*(N-1) = -\underbrace{(R + B^T P(0)B)^{-1} B^T P(0)A}_{\triangleq K(N-1)} \mathbf{x}(N-1)$$

$$\triangleq K(N-1)$$

Discrete-Time LQR – Backward Recursion Step N-2

Note: $J_{\mathbf{x}(N-1) \rightarrow \mathbf{x}(N)}^* = \mathbf{x}^T(N-1)P(1)\mathbf{x}(N-1)$ where:

$$P(1) = (A - BK(N-1))^T P(0)(A - BK(N-1)) + K^T(N-1)RK(N-1) + Q$$

Continuing backward recursion (step N-1):

$$\begin{aligned} J_{\mathbf{x}(N-2) \rightarrow \mathbf{x}(N)} &= \mathbf{x}^T(N-2)Q\mathbf{x}(N-2) + Ru^2(N-2) + \mathbf{x}^T(N-1)P(1)\mathbf{x}(N-1) \\ &= \mathbf{x}^T(N-2)Q\mathbf{x}(N-2) + Ru^2(N-2) + (A\mathbf{x}(N-2) + Bu(N-2))^T P(1)(A\mathbf{x}(N-2) + Bu(N-2)) \end{aligned}$$

Key observation: The cost to go above is **identical** in structure to $J_{\mathbf{x}(N-1) \rightarrow \mathbf{x}(N)}$, except that:

- $u(N-1)$ and $\mathbf{x}(N-1)$ have been replaced with $u(N-2)$ and $\mathbf{x}(N-2)$
- $P(0)$ has been replaced with $P(1)$

Result: The optimal control signal ($u^*(N-2)$) will be the same as before, except that $\mathbf{x}(N-1)$ and $P(0)$ will be replaced with $\mathbf{x}(N-2)$ and $P(1)$, respectively

Discrete-Time LQR – Derivation in Class

Given: $\underline{x}(k+1) = A\underline{x}(k) + B u(k)$

Minimize $J(\underline{u}; \underline{x}(0)) = \underline{x}^T(N) S \underline{x}(N) + \sum_{i=0}^{N-1} (\underline{x}^T(i) Q \underline{x}(i) + R u^2(i))$

Equivalent to minimizing

$$J(\underline{u}; \underline{x}(0)) = \sum_{i=0}^{N-1} [\underline{x}^T(i+1) Q \underline{x}(i+1) + R u^2(i)] \quad \text{when } S = Q$$

$$= \underline{x}^T(N) Q \underline{x}(N) + \sum_{i=0}^{N-1} [\underline{x}^T(i) Q \underline{x}(i) + R u^2(i)] - \underbrace{\underline{x}^T(0) Q \underline{x}(0)}_{\text{Constant}}$$

Constant
 \Rightarrow no impact
 on minimizer

$P(0) \triangleq S$

$$J_{\underline{x}(N) \rightarrow \underline{x}(N)}^* = \underline{x}^T(N) P(0) \underline{x}(N)$$

Just as the purpose (and Iraq war metaphor) of the strange agrarian community in M. Night Shyamalan's *The Village* is initially hazy but becomes remarkably clear later on, the purpose of the mysterious assignment of $P(0)$ will become clear...with stunning implications for LQR control

Discrete-Time LQR – Derivation in Class

Stage $N-1$:

$$\begin{aligned}
 J_{\underline{x}(N-1) \rightarrow \underline{x}(N)} &= \underline{x}^T(N-1)Q\underline{x}(N-1) + Ru^2(N-1) + \underline{x}^T(N)P(0)\underline{x}(N) \\
 &= \underline{x}^T(N-1)Q\underline{x}(N-1) + Ru^2(N-1) + [A\underline{x}(N-1) + Bu(N-1)]^T P(0) \dots \\
 &\quad [A\underline{x}(N-1) + Bu(N-1)]
 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial J_{\underline{x}(N-1) \rightarrow \underline{x}(N)}}{\partial u(N-1)} &= 2Ru(N-1) + 2B^T P(0)Bu(N-1) \\
 &\quad + 2B^T P(0)A\underline{x}(N-1) = 0 \text{ when} \\
 &\quad u(N-1) = u^*(N-1)
 \end{aligned}$$

$$\Rightarrow u^*(N-1) = - \underbrace{(R + B^T P(0)B)^{-1} (B^T P(0)A)}_{K(N-1)} \underline{x}(N-1)$$

Discrete-Time LQR – Derivation in Class

$$\begin{aligned}
 J_{\underline{x}(N-1) \rightarrow \underline{x}(N)}^* &= \underline{x}^T(N-1) Q \underline{x}(N-1) + \underline{x}^T(N-1) K^T(N-1) \overset{R}{K}(N-1) \underline{x}(N-1) \\
 &\quad + \left[(A - BK(N-1)) \underline{x}(N-1) \right]^T P(0) \left[(A - BK(N-1)) \underline{x}(N-1) \right] \\
 &= \underline{x}^T(N-1) \underbrace{\left[Q + K^T(N-1) R K(N-1) + (A - BK(N-1))^T P(0) (A - BK(N-1)) \right]}_{P(1)} \underline{x}(N-1)
 \end{aligned}$$

$$\Rightarrow J_{\underline{x}(N-1) \rightarrow \underline{x}(N)}^* = \underline{x}^T(N-1) P(1) \underline{x}(N-1)$$

Same as at stage N, with $P(0)$ replaced by $P(1)$ (ahh...so that's the reason for $P(i)$) and $\underline{x}(N)$ replaced with $\underline{x}(N-1)$

...At stage N-2, we have an identical optimization problem to the one at stage N-1, with $\underline{x}(N-1)$ replaced with $\underline{x}(N-2)$, $u(N-1)$ replaced with $u(N-2)$, and $P(0)$ replaced with $P(1)$

...And if the problems are identical (with index substitutions), then the solutions will be identical (with index substitutions)!

Discrete-Time LQR – General Control Law

Remember: $u^*(N-1) = -K(N-1)\mathbf{x}(N-1)$ and $u^*(N-2) = -K(N-2)\mathbf{x}(N-2)$

where: $K(N-1) = -(R + B^T P(0)B)^{-1} B^T P(0)A$ and $K(N-2) = -(R + B^T P(1)B)^{-1} B^T P(1)A$

where: $P(0) = S$ and $P(1) = (A - BK(N-1))^T P(0)(A - BK(N-1)) + K^T(N-1)RK(N-1) + Q$

Continuing the backward recursion (generic step N-i):

$$u^*(N-i) = -K(N-i)\mathbf{x}(N-i)$$

where: $K(N-i) = -(R + B^T P(i-1)B)^{-1} B^T P(i-1)A$

where: $P(i) = (A - BK(N-i))^T P(i-1)(A - BK(N-i)) + K^T(N-i)RK(N-i) + Q$

Key takeaway: The resulting controller is a linear, time-varying controller whose gains at each step can be found recursively!

Discrete-Time LQR – Block Diagram and Observations

Key takeaway (reminder): The resulting controller is a linear, time-varying controller whose gains at each step can be found recursively!

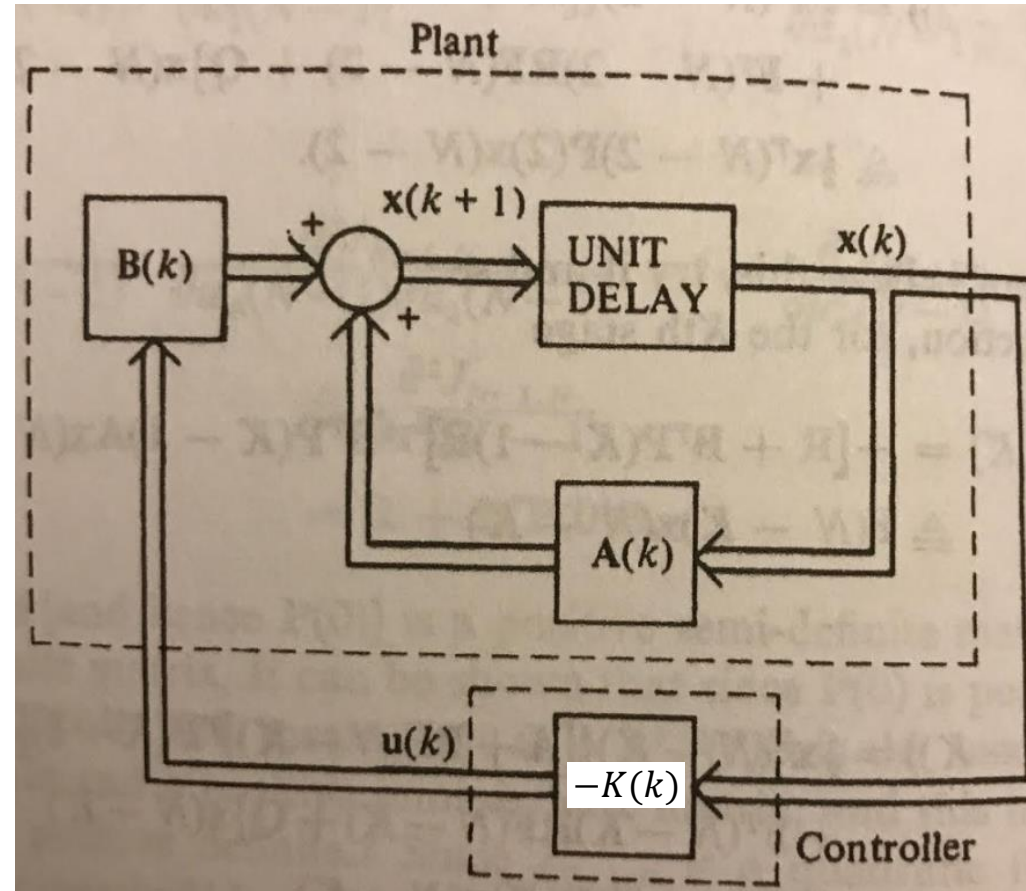


Image source:
Kirk, page 82

Key differences from standard DP:

- No quantization of control and states (since cost can be differentiated)
- Unlike DP, where ***a sequence of control signals is computed all at once***, LQR lends itself to a ***feedback control law*** (which happens to be linear and time-varying)

Infinite Horizon LQR - Setup

Consider the following system dynamics:

$$\mathbf{x}(k + 1) = A\mathbf{x}(k) + Bu(k) \quad \mathbf{x} \in \mathbb{R}^n, u \in \mathbb{R}$$

...and the following objective function:

$$J(\mathbf{x}(0), \mathbf{u}) = \sum_{i=0}^{\infty} (\mathbf{x}^T(i)Q\mathbf{x}(i) + Ru^2(i))$$

where: $Q \succ 0, R > 0$

...Same problem as before, except that:

- $N = \infty$
- There is no terminal penalty

Infinite Horizon LQR - Solution

Key facts:

- In the finite-horizon LQR problem, $J_{\mathbf{x}(i) \rightarrow \mathbf{x}(N)}^* = \mathbf{x}^T(i)P(i)\mathbf{x}(i)$, regardless of the presence of a terminal penalty
- For a time-invariant system, the infinite horizon cost to go from a given state (\mathbf{x}) is the same regardless of the time instant at which the controller is “turned on”:

$$J_{\mathbf{x}(i) \rightarrow \infty}^* = J_{\mathbf{x}(i+1) \rightarrow \infty}^*$$

- Putting these facts together, it must be true that $P(i) = P(i+1) = P \dots P$ is a constant!

Recall the update laws for P and K : $K(N-i) = -(R + B^T P(i-1)B)^{-1} B^T P(i-1)A$

$$P(i) = (A - BK(N-i))^T P(i-1)(A - BK(N-i)) + K^T(N-i)RK(N-i) + Q$$

The “update” laws now become: $K = -(R + B^T P B)^{-1} B^T P A$

$$P = (A - BK)^T P (A - BK) + K^T R K + Q$$

Discrete-time algebraic
Riccati equation

Infinite Horizon LQR - Solution

Remember: $J_{\underline{x}(i) \rightarrow \underline{x}(\infty)}^* = J_{\underline{x}(i+1) \rightarrow \underline{x}(\infty)}^*$
Same state values

$$\Rightarrow J_{\underline{x}(i) \rightarrow \underline{x}(\infty)}^* = \underline{x}^T(i) P(\infty-i) \underline{x}(i)$$

$$J_{\underline{x}(i+1) \rightarrow \underline{x}(\infty)}^* = \underline{x}^T(i+1) P(\infty-i-1) \underline{x}(i+1)$$

$$\text{If } \underline{x}(i) = \underline{x}(i+1), P(\infty-i-1) = P(\infty-i)$$

Infinite Horizon LQR - Solution

Given $\bar{x}^{(k+1)} = A\bar{x}^{(k)} + Bu^{(k)}$, we can design a state feedback - based regulator in 2 ways:

1) Pole placement:

$$u = -K\bar{x}(k)$$

$$\Rightarrow \bar{x}_{(k+1)} = \underbrace{(A - BK)}_{\text{Set eigenvalues equal to desired pole locations.}} \bar{x}(k)$$

Set eigenvalues equal to desired pole locations.

2) LQR design:

Specify Q, R ($Q \in \mathbb{R}^{N \times N}$, $R \in \mathbb{R}$)

$$\text{where } J(\bar{x}_0, K) = \sum_{i=0}^{\infty} \{ \bar{x}^T(i) Q \bar{x}(i) + R u(i)^2 \}$$

\Rightarrow Determine K based on discrete-time algebraic Riccati equation.

Infinite Horizon LQR Solution in MATLAB



Syntax:

`K = dlqr(A,B,Q,R);`

Resulting feedback gain vector

Example to work in class:

$$\mathbf{x}(k+1) = A\mathbf{x}(k) + Bu(k)$$

$$A = \begin{bmatrix} 0 & 1 \\ -2 & -3 \end{bmatrix} \quad B = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$J(\mathbf{x}(0), \mathbf{u}) = \sum_{i=0}^{\infty} (\mathbf{x}^T(i)Q\mathbf{x}(i) + Ru^2(i))$$

$$Q = \begin{bmatrix} 2 & 0 \\ 0 & 1 \end{bmatrix} \quad R = 1$$

Observations About the LQR Solution



Good features:

- Closed-form solution available (via differentiation of cost to go), so no need to quantize states and control variables
- Optimal controller is linear state feedback...can be implemented as a feedback controller, rather than just a sequence of control values that begins at time 0
- With an infinite horizon and time-invariant system, the feedback gain is time-invariant

Not so good features:

- Quadratic cost function (in states and control variables) required
- No constraints considered

Question to be addressed in future lectures: Can we obtain the good features above while allowing for more general cost function structures and constraints?

Preview of Upcoming Lectures



Model predictive control – answers the question of “how do we take the trajectory optimization algorithms that we’ve learned about and use them for *real-time feedback control*?”:

- Basic MPC setup and notation (October 19)
- MPC implementation in MATLAB and Simulink (October 24)
- Stability and robustness of MPC (October 26)